

```

In [4]: # for numerical computing
import numpy as np
# for dataframes
import pandas as pd
# for easier visualization
import seaborn as sns
# for visualization and to display plots
from matplotlib import pyplot as plt
%matplotlib inline
# import color maps
from matplotlib.colors import ListedColormap
# Ignore Warnings
import warnings
warnings.filterwarnings("ignore")
from math import sqrt
# to split train and test set
from sklearn.model_selection import train_test_split
# to perform hyperparameter tuning
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.linear_model import Ridge # Linear Regression + L2 regularizat
from sklearn.linear_model import Lasso # Linear Regression + L1 regularizat
from sklearn.svm import SVR # Support Vector Regressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
# Evaluation Metrics
from sklearn.metrics import mean_squared_error as mse
from sklearn.metrics import r2_score as rs
from sklearn.metrics import mean_absolute_error as mae
#import xgboost
import os
mingw_path = 'C:\\Program Files\\mingw-w64\\x86_64-7.2.0-posix-seh-rt_v5-rev0\\m
os.environ['PATH'] = mingw_path + ';' + os.environ['PATH']
# to save the final model on disk
from sklearn.externals import joblib

```

Loading Black Friday Data

```

In [5]: df = pd.read_csv('BlackFriday 2.csv')

```

```

In [6]: df.shape

```

```

Out[6]: (537577, 12)

```

In [7]: `df.columns`

Out[7]: Index(['User_ID', 'Product_ID', 'Gender', 'Age', 'Occupation', 'City_Category', 'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category_1', 'Product_Category_2', 'Product_Category_3', 'Purchase'], dtype='object')

In [8]: `df.head()`

Out[8]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Mar
0	1000001	P00069042	F	0-17	10	A	2	
1	1000001	P00248942	F	0-17	10	A	2	
2	1000001	P00087842	F	0-17	10	A	2	
3	1000001	P00085442	F	0-17	10	A	2	
4	1000002	P00285442	M	55+	16	C	4+	

Filtering the categorical data

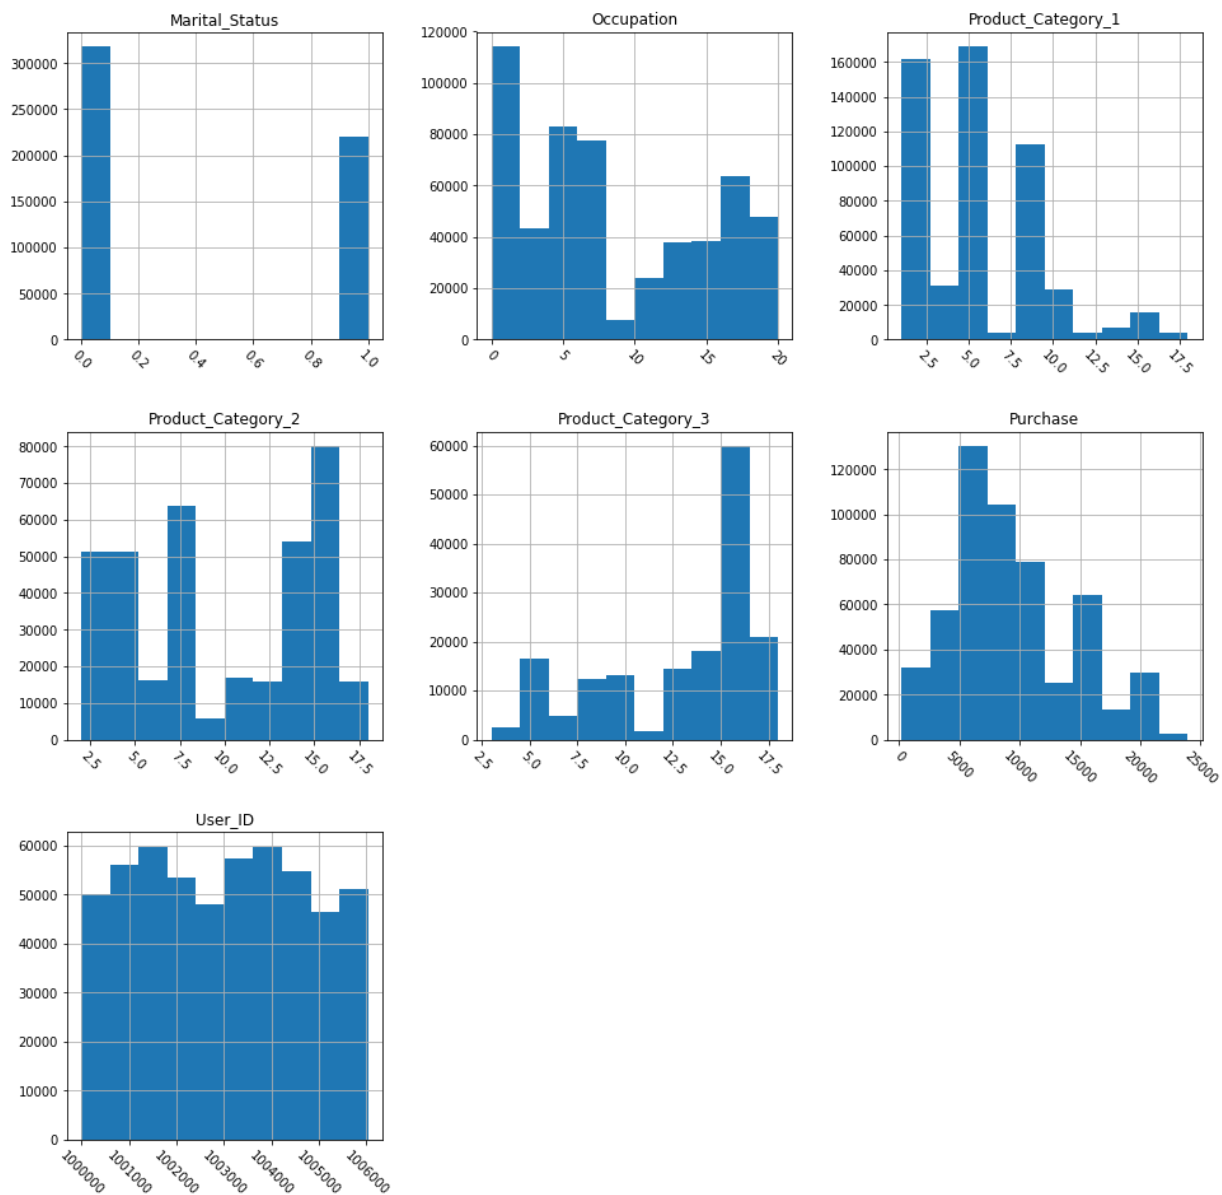
In [9]: `df.dtypes[df.dtypes == 'object']`

Out[9]:

Product_ID	object
Gender	object
Age	object
City_Category	object
Stay_In_Current_City_Years	object
dtype:	object

Distribution of numerical data

```
In [10]: df.hist(figsize=(16,16), xrot=-45)  
plt.show()
```



```
In [11]: df.describe()
```

```
Out[11]:
```

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3
count	5.375770e+05	537577.000000	537577.000000	537577.000000	370591.000000	169916.000000
mean	1.002992e+06	8.08271	0.408797	5.295546	9.842144	14.451162
std	1.714393e+03	6.52412	0.491612	3.750701	5.087259	6.359382
min	1.000001e+06	0.00000	0.000000	1.000000	2.000000	3.000000
25%	1.001495e+06	2.00000	0.000000	1.000000	5.000000	8.000000
50%	1.003031e+06	7.00000	0.000000	5.000000	9.000000	13.000000
75%	1.004417e+06	14.00000	1.000000	8.000000	15.000000	17.000000
max	1.006040e+06	20.00000	1.000000	18.000000	18.000000	18.000000



Distribution of categorical data

```
In [12]: df.describe(include=['object'])
```


```
Out[12]:
```

	Product_ID	Gender	Age	City_Category	Stay_In_Current_City_Years
count	537577	537577	537577	537577	537577
unique	3623	2	7	3	5
top	P00265242	M	26-35	B	1
freq	1858	405380	214690	226493	189192

Bar plots for categorical data


```
In [13]: plt.figure(figsize = (8,8))  
sns.countplot(y='Age',data = df)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x25033c3f780>
```



```
In [14]: plt.figure(figsize = (8,8))  
sns.countplot(y='Stay_In_Current_City_Years',data = df)
```


```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x25034204748>
```



Segmentations

```
In [15]: sns.boxplot(y='Gender', x='Purchase', data = df)
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x25033f462e8>
```



Comparing two genders across other features


```
In [16]: df.groupby('Gender').mean()
```

```
Out[16]:
```

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3
Gender						
F	1.003088e+06	6.742672	0.417733	5.595445	10.007969	0.000000
M	1.002961e+06	8.519705	0.405883	5.197748	9.789072	0.000000

```
In [17]: sns.boxplot(y='Gender', x='Occupation', data = df)
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x25033cb1d68>
```



```
In [18]: df.groupby('Gender').agg([np.mean,np.std])
```

```
Out[18]:
```

	User_ID		Occupation		Marital_Status		Product_Category_1		r
	mean	std	mean	std	mean	std	mean	std	
Gender									
F	1.003088e+06	1774.236455	6.742672	6.242116	0.417733	0.493188	5.595445	3.476495	1
M	1.002961e+06	1693.251916	8.519705	6.554518	0.405883	0.491063	5.197748	3.830816	

```
In [19]: plt.figure(figsize=(20,20))
```

```
Out[19]: <Figure size 1440x1440 with 0 Axes>
```

```
<Figure size 1440x1440 with 0 Axes>
```


```
In [20]: df.corr()
```

```
Out[20]:
```

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase
User_ID	1.000000	-0.023024	0.018732	0.003687	0.001471	0.004045	0.005389
Occupation	-0.023024	1.000000	0.024691	-0.008114	-0.000031	0.013452	0.021104
Marital_Status	0.018732	0.024691	1.000000	0.020546	0.015116	0.019452	0.000129
Product_Category_1	0.003687	-0.008114	0.020546	1.000000	0.540423	0.229490	-0.314125
Product_Category_2	0.001471	-0.000031	0.015116	0.540423	1.000000	0.543541	-0.209971
Product_Category_3	0.004045	0.013452	0.019452	0.229490	0.543541	1.000000	0.000000
Purchase	0.005389	0.021104	0.000129	-0.314125	-0.209971	0.000000	1.000000


```
In [21]: sns.heatmap(df.corr())
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x2503433ecf8>
```




```
In [22]: corr = df.corr()
print (corr['Purchase'].sort_values(ascending=False)[:4], '\n')
print ('-----')
print (corr['Purchase'].sort_values(ascending=False)[-3:])`
```

```
Purchase          1.000000
Occupation        0.021104
User_ID           0.005389
Marital_Status    0.000129
Name: Purchase, dtype: float64
```

```
-----
Product_Category_3 -0.022257
Product_Category_2 -0.209973
Product_Category_1 -0.314125
Name: Purchase, dtype: float64
```

```
In [23]: from matplotlib.colors import ListedColormap
mask = np.zeros_like(df.corr())
mask[np.triu_indices_from(mask)] = True
plt.figure(figsize = (10,10))
with sns.axes_style("white"):
    ax = sns.heatmap(df.corr()*100,mask =mask, fmt = '.0f',
                    annot = True, lw=1,cmap =ListedColormap(["green","yellow","red","blue"]))
```



Data Cleaning

Drop any duplicate

```
In [24]: df = df.drop_duplicates()
df.shape
```

```
Out[24]: (537577, 12)
```

```
In [25]: df.City_Category.unique()
```

```
Out[25]: array(['A', 'C', 'B'], dtype=object)
```

```
In [26]: df.Stay_In_Current_City_Years.unique()
```

```
Out[26]: array(['2', '4+', '3', '1', '0'], dtype=object)
```

```
In [27]: df.Product_Category_2.unique()
```

```
Out[27]: array([nan,  6., 14.,  2.,  8., 15., 16., 11.,  5.,  3.,  4., 12.,  9.,
                10., 17., 13.,  7., 18.])
```

```
In [28]: df.Product_Category_2.fillna(9, inplace = True)
```

```
In [29]: df.Product_Category_2.unique()
```


```
Out[29]: array([ 9.,  6., 14.,  2.,  8., 15., 16., 11.,  5.,  3.,  4., 12., 10.,
                17., 13.,  7., 18.])
```

Removing Outliers

```
In [30]: # Outliers can cause problems with certain types of models.
# Boxplots are a nice way to detect outliers
# Let's start with a box plot of your target variable, since that's what you're c
```


```
In [31]: sns.boxplot(df.Purchase)
```

```
Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x25033e990b8>
```




```
In [32]: sns.boxplot(df.Product_Category_1)
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x25033df9f98>
```




```
In [33]: sns.boxplot(df.Product_Category_2)
```

```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0x25034414cf8>
```



```
In [34]: sns.boxplot(df.Product_Category_3)
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x2503447f358>
```



Label missing categorical data

```
In [35]: # Display number of missing values by categorical feature
df.select_dtypes(include=['object']).isnull().sum()
```

```
Out[35]: Product_ID          0
Gender          0
Age            0
City_Category   0
Stay_In_Current_City_Years  0
dtype: int64
```

Flag and Fill missing numerical data

```
In [36]: # Display number of missing values by numeric feature
df.select_dtypes(exclude=['object']).isnull().sum()
```

```
Out[36]: User_ID          0
Occupation    0
Marital_Status  0
Product_Category_1  0
Product_Category_2  0
Product_Category_3  373299
Purchase      0
dtype: int64
```

```
In [37]: df['Product_Category_3'] = df['Product_Category_3'].fillna(df['Product_Category_3'].select_dtypes(exclude=['object']).isnull().sum())
```

```
Out[37]: User_ID      0
Occupation  0
Marital_Status  0
Product_Category_1  0
Product_Category_2  0
Product_Category_3  0
Purchase      0
dtype: int64
```

```
In [39]: # Save cleaned dataframe to new file
# "C:\Users\Shakena Ford\Desktop\cleaneddf.csv"
df.to_csv(r'C:\Users\Shakena Ford\Desktop\cleaned.csv', index=False)
```

Encode Dummy Variables

```
In [40]: # Machine Learning algorithms cannot directly handle categorical features. Specifically,
# Therefore, we need to create dummy variables for our categorical features.
# Dummy variables are a set of binary (0 or 1) features that each represent a single category.
# Create a new dataframe with dummy variables for our categorical features.
```

```
In [41]: df = pd.get_dummies(df, columns=['Gender', 'Age', 'City_Category'])
```


```
In [42]: # Note: There are many ways to perform one-hot encoding,
# you can also use LabelEncoder and OneHotEncoder classes in SKLEARN or use the get_dummies method
```

```
In [43]: df.head()
```

```
Out[43]:
```

	User_ID	Product_ID	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3
0	1000001	P00069042	10	2	0	0	0	0
1	1000001	P00248942	10	2	0	0	0	0
2	1000001	P00087842	10	2	0	0	0	1
3	1000001	P00085442	10	2	0	0	0	1
4	1000002	P00285442	16	4+	0	0	0	0

5 rows × 21 columns



```
In [45]: # Save cleaned dataframe to new file
# "C:\Users\Shakena Ford\Desktop\cleaneddf.csv"
df.to_csv(r'C:\Users\Shakena Ford\Desktop\analytical.csv', index=False)
```

Machine Learning

Data Preparation

```
In [48]: df = pd.read_csv("analytical.csv")
```

```
In [49]: y = df.Purchase
x = df.drop('Purchase', axis = 1)
```

```
In [50]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_
```

```
In [51]: print(x_train.shape, x_test.shape, y_train.shape, y_test.shape)

(430061, 20) (107516, 20) (430061,) (107516,)
```

Data Standardization

```
In [*]:
```

```
In [*]: train_mean = x_train.mean()
train_std = x_train.std()
```

```
In [*]: x_train = (x_train - train_mean) / train_std
```

```
In [*]: x_train.describe()
```

```
In [*]: x_test = (x_test - train_mean) / train_std
```

```
In [*]: x_test.describe()
```

Baseline Model

In [*]:

In [*]:

```
y_train_pred = np.ones(y_train.shape[0])*y_train.mean()
```

In [*]:

```
## Predict Test results  
y_pred = np.ones(y_test.shape[0])*y_train.mean()  
from sklearn.metrics import r2_score
```

In [*]:

```
print("Train Results for Baseline Model:")  
print("*****")  
print("Root mean squared error: ", sqrt(mse(y_train.values, y_train_pred)))  
print("R-squared: ", r2_score(y_train.values, y_train_pred))  
print("Mean Absolute Error: ", mae(y_train.values, y_train_pred))
```

In []:

In []: